

Alexey Grigorev

Coding agents produce data

Token usage, models, tool calls, cache, and project traces are all useful signals.

The traces are already there

Local coding agents store metadata about sessions. This can include usage, models, inputs, outputs, cache, and project-level activity.

Raw JSON is a poor dashboard

You can open trace files by hand, but it quickly becomes painful. The useful questions need aggregation, filtering, and repeatable queries.

The data workflow is simple

Read the trace files, load them with dlt, store them in DuckDB, and build a dashboard for usage by day, model, and project.

This is agent observability

Start with traces as data. Then add ingestion, schema, validation, and reports around them.

Watch the workshop

Ingesting Agent Traces with dlthub, with Alena Astrakhantseva.