

Alexey Grigorev

Asking an agent to re-scan logs does not scale

A one-off question is fine. Repeated analytics needs a pipeline.

The one-off prompt works

You can ask an agent: how many tokens did I use last month? It can read local trace files, aggregate them, and answer.

The next question repeats work

Change the time range or grouping and the agent may need to read and aggregate the same raw logs again. This spends time and tokens.

Make it reusable

Ingest the traces once, normalize the structure, store them in a queryable destination, and refresh on a schedule.

Then it becomes analytics

Token usage, model usage, project activity, and cost patterns become normal dashboard questions.

Watch the workshop

Alena Astrakhantseva shows the workflow with dlt and agent traces.