

Alexey Grigorev

Agent traces rarely arrive as clean tables

They arrive as nested JSON. Schema design becomes part of observability.

The first pipeline created 78 tables

dlt normalized the trace files into relational tables. The high table count came from deeply nested messages, usage, cache, models, and tool data.

This is expected

Nested JSON has structure. Loading it into a relational destination requires choices about how much to flatten.

Schema pollution is real

Flatten everything and the schema can explode.
Keep too much raw JSON and important fields
become harder to query.

The adjusted schema had 40 tables

The agent kept selected fields as JSON columns. The result was a more manageable schema while preserving useful structure.

Watch the workshop

A practical data engineering lesson inside an agent traces demo.